

Lecture 5: PIR \rightarrow OT and Preprocessing PIR

Scribe: Tanisha Saxena

February 14, 2024

1 PIANO: An Extremely Simple PIR

1.1 Motivation

Recall that in previous lectures, we talked about classical PIR where the server stores the original database (DB) unencoded and there's no preprocessing. Classical PIR requires the server to look through each element of the database, otherwise the server can deduce that the skipped elements are unimportant to the client. Beimel et. al [BIM00] proved that a server in classical PIR must have linear computation per query. Linear computation will unlikely scale for many real-world applications, e.g., private DNS and private web search. To avoid this linear computation barrier, more recent PIR schemes considered the preprocessing model [BIM00]. [elaine: cite also corrigan-gibbs and kogan] In this lecture, we will show how we can achieve sublinear computation per query in the preprocessing model.

1.2 Background

Previous works have considered main models of pre-processing:

- **Public preprocessing:** The server computes an encoding of the database DB. This preprocessing is shared globally by all clients.
- **Client-specific preprocessing:** The server performs a preprocessing protocol with each client. The client is stateful, i.e., each client maintains some local state called the **hint**.

The “doubly-efficient PIR” scheme Wei-Kai talked about in his guest lecture is in the public preprocessing model. In this lecture, we will describe a scheme called Piano by Zhou et al. [ZPSZ23] that uses the client-specific preprocessing model.


1.3 Goals

Piano achieves $O(\sqrt{n})$ time per query and $O(\sqrt{n})$ bandwidth [ZPSZ23] assuming $\tilde{O}(\sqrt{n})$ client space. The scheme assumes that pseudorandom functions (PRF) exists (One-way Function).

For the sake of our first pass, we're going to assume each query is random (i.e. the index of each message the client wants is randomly chosen), and that there are up to \sqrt{n} unique queries. Basically, we're assuming our current queries are bounded and random. We'll remove this restriction later and are only using it now for the sake of ease.

1.4 The Scheme

Given a database DB, we split the n values into \sqrt{n} chunks of size \sqrt{n} .



And as easily as that, we have removed the bounded and random queries restrictions!

Applications This scheme is much faster than non-processing PIR schemes. Because of that, it has many applications for fast secrecy. For example, `haveibeenpwned.com`, a website that checks if a user’s password has been in a data breach, would be able to more efficiently run a search on a user’s password without actually ever knowing what the password is.

References

- [BIM00] Amos Beimel, Yuval Ishai, and Tal Malkin. Reducing the servers computation in private information retrieval: Pir with preprocessing. In *Advances in Cryptology—CRYPTO 2000: 20th Annual International Cryptology Conference Santa Barbara, California, USA, August 20–24, 2000 Proceedings 20*, pages 55–73. Springer, 2000.
- [ZPSZ23] Mingxun Zhou, Andrew Park, Elaine Shi, and Wenting Zheng. Piano: Extremely simple, single-server pir with sublinear server computation. *Cryptology ePrint Archive*, 2023.